

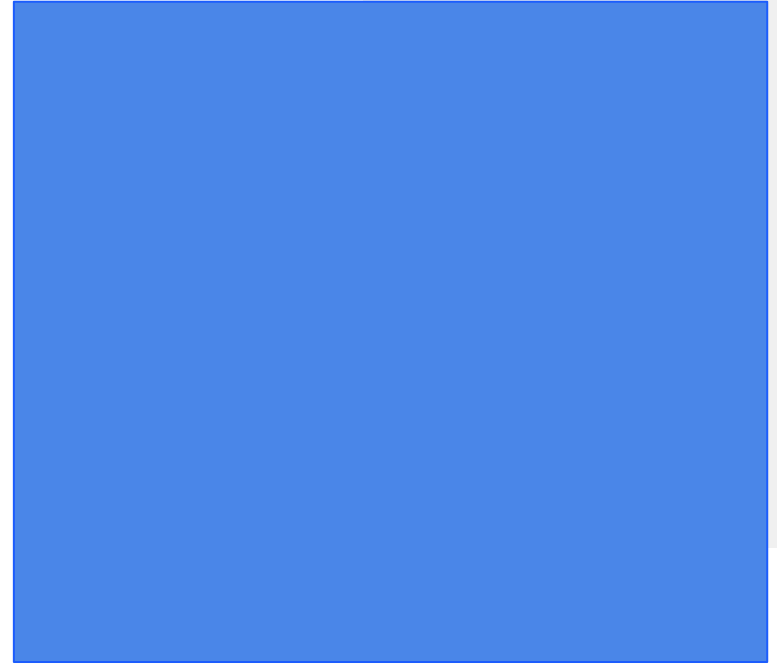


Big Data Analytics

Dr Sirintra Vaiwsri | Email: sirintra.v@itm.kmutnb.ac.th



Structured API Execution






Steps (Chambers and Zaharia, 2018)



1. Write DataFrame/SQL Code
2. Spark converts code to a Logical Plan if the code is valid.
3. Spark transforms the Logical Plan to a Physical Plan.
4. Spark executes the Physical Plan on a cluster.



Logical Plan (Chambers and Zaharia, 2018)

1. The code is first converted to an unresolved logical plan.
 2. Spark then uses a “catalogue” to resolve tables and columns in the analyser.
 3. The analyser rejects an unresolved logical plan if the referred tables do not exist, otherwise, it accepts which becomes a resolved logical plan.
 4. Spark passes the resolved logical plan through the catalyst optimiser to optimise the resolved logical plan resulting in the “optimised logical plan”
- 

Physical Plan (Chambers and Zaharia, 2018)

- Physical plan, also called Spark plan, generates different physical executions and compares them through a cost model.
- It can manage the execution of an optimised logical plan on a cluster.
- It results in RDDs and their transformations, in other words, the Physical Plan runs the code over the RDDs.

Dataframes



Dataframes


(Antolínez García, 2023; Chambers and Zaharia, 2018)

- Dataframes allow Spark to efficiently conduct shuffling by moving data across nodes.
- It represents a table of data containing rows and columns.
- Schema includes column name, data type, and nullable flag.

```
root
|-- status_id: string (nullable = true)
|-- status_type: string (nullable = true)
|-- status_published: string (nullable = true)
|-- num_reactions: string (nullable = true)
|-- num_comments: string (nullable = true)
|-- num_shares: string (nullable = true)
|-- num_likes: string (nullable = true)
|-- num_loves: string (nullable = true)
|-- num_wows: string (nullable = true)
|-- num_hahas: string (nullable = true)
|-- num_sads: string (nullable = true)
|-- num_angrys: string (nullable = true)
```



Dataframes (Antolínez García, 2023)

- Dataframes can be distributed across nodes to support distributed computing architecture.
 - Users can query using SQL which is then sent to Dataframes and managed by the Catalyst Optimiser.
 - Catalyst Optimiser is responsible for building query execution.
 - Dataframes can be created from external data sources such as CSV and JSON files.
- 

Dataframes

```
from pyspark.sql import SparkSession
# Create Spark Session
read_file = spark.read.format\
    (<FILE FORMAT>).option("header",<True/False>).\
    load(<FILE NAME>) # Read file
read_file.printSchema() # Print schema
```

```
root
|-- status_id: string (nullable = true)
|-- status_type: string (nullable = true)
|-- status_published: string (nullable = true)
|-- num_reactions: string (nullable = true)
|-- num_comments: string (nullable = true)
|-- num_shares: string (nullable = true)
|-- num_likes: string (nullable = true)
|-- num_loves: string (nullable = true)
|-- num_wows: string (nullable = true)
|-- num_hahas: string (nullable = true)
|-- num_sads: string (nullable = true)
|-- num_angrys: string (nullable = true)
```

Spark SQL (Antolínez García, 2023)

- Spark SQL can query data directly from file.

```
sqlDF = read_File.select(read_file.<column1>,\n    read_file.<column2>) # SQL query\nsqlDF.show(<Row Numbers>) # Show result
```

Spark SQL (Antolínez García, 2023)

- Spark SQL query data directly from file with **where** clause.

```
from pyspark.sql.types import *  
sqlDF = read_File.select(read_file.<column name1>,\  
    read_file.<column name 2>).\  
    where(read_file.num_reactions.\  
        cast(IntegerType()) > <some numbers>).\  
    withColumnRenamed(<old name>, \  
        <new name>).\  
    orderBy(readfile.<column name>) # SQL query  
sqlDF.show(<Row Numbers>) # Show result
```

status_id	Reactions
246675545449582_7...	3115
246675545449582_7...	3190
246675545449582_7...	3510

only showing top 3 rows

Spark SQL (Antolínez García, 2023)

- Spark SQL can conduct directly over the file through temporary views by using `createOrReplaceTempView(<temp name>)`.

```
read_file.createOrReplaceTempView(<temp name>) # Create
                                              # temporary view
sqlDF = spark.sql("select * from <temp name>") # SQL query
sqlDF.show(<Row Numbers>) # Show result
```

status_id	status_type	status_published	num_reactions	num_comments	num_shares	num_likes	num_loves	num_wows	num_hahas	num_sads	num_angrys
246675545449582_1...	video	4/22/2018 6:00	529	512	262	432	92	3	1	1	0
246675545449582_1...	photo	4/21/2018 22:45	150	0	0	150	0	0	0	0	0
246675545449582_1...	video	4/21/2018 6:17	227	236	57	204	21	1	1	0	0

only showing top 3 rows



Save Modes (Antolínez García, 2023)

- Data can be saved into a file.
- Save Modes are:
 - errorifexists or error - exception is sent if data already exists
 - append - data is appended to the destination
 - overwrite - data is overwritten if data already exists
 - ignore - data will be ignored if data already exists

```
sqlDF.write.mode("overwrite").csv("<folder>") # Save data
```



Random Split (Chambers and Zaharia, 2018)

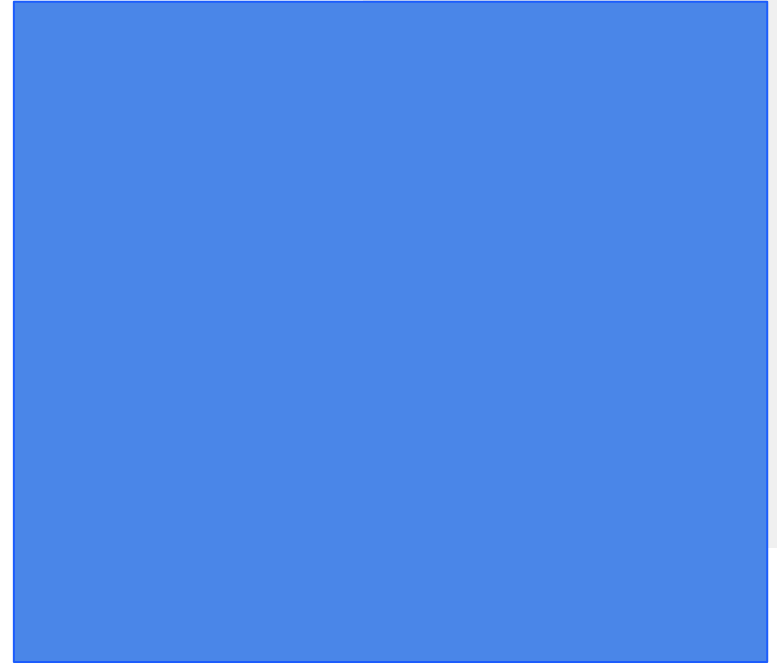
- randomSplit is used for splitting data in a Dataframe.
- It is useful when creating training and test sets are required such as when working with machine learning.

```
split = sqlDF.randomSplit([number1, number2]) # Split data  
where numbers 1 and 2 are sum to 1.0.
```

```
split[0].show() # Show data in the first set (number1)
```

```
split[1].show() # Show data in the second set (number2)
```

Aggregations



Aggregations (Chambers and Zaharia, 2018)

- Aggregation collects data together and produces one result for each group.
- Aggregations are available as functions such as count, first and last, min and max, sum, etc.

fb_live_thailand2 and
fb_live_thailand3

```
read_file = spark.read.format("csv").\  
    option("header", True).\  
    load("data/*.csv") # Multiple files  
  
# Create a temporary view  
# Select all from the temporary view  
# Show result
```


Aggregation Functions (Chambers and Zaharia, 2018)

- count is used for counting rows.

```
from pyspark.sql.functions import count
```

```
# Select all from the temporary view
```

```
sqlDF.select(count("<column name>")) # Select count
```

```
sqlDF = spark.sql("select * from tempTable where \  
    tempTable.status_type == 'photo'") # Select...where
```

```
sqlDF.select(count("status_published")) # Select count
```

```
+-----+  
|count(status_published)|  
+-----+  
|                        22|  
+-----+
```

Aggregation Functions (Chambers and Zaharia, 2018)

- countDistinct is used for counting number of unique groups.

```
from pyspark.sql.functions import countDistinct
# Select all from the temporary view
sqlDF.select(countDistinct("<column name>")) # Select
                                             #countDistinct
```

```
+-----+
|count(DISTINCT status_type)|
+-----+
|                           2|
+-----+
```

Aggregation Functions (Chambers and Zaharia, 2018)

- first and last are used to retrieve the first and last rows from a Dataframe.

```
from pyspark.sql.functions import first, last
# Select all from the temporary view
sqlDF.select(first("<column name>"), \
              last("<column name>")) # Select first and last
```

first(status_published)	last(status_published)
4/22/2018 6:00	3/23/2018 7:09

Aggregation Functions (Chambers and Zaharia, 2018)

- min and max are used for finding the minimum and maximum values from a Dataframe.

```
from pyspark.sql.functions import min, max
from pyspark.sql.types import *
```

```
# Select all from the temporary view
```

```
sqlDF = sqlDF.withColumn('<new column name>', \
    sqlDF['<old column name>']).\
```

```
    cast(IntegerType())) # Convert the column to integer
                           # then add as a new column name
```

```
sqlDF.select(min("<new column name>"), \
    max("<new column name>")) # Select min and max
```

min(reactions)	max(reactions)
18	529

Aggregation Functions (Chambers and Zaharia, 2018)

- sum and sumDistinct are used for summing a total where sumDistinct sums a distinct set of values.

```
from pyspark.sql.functions import sum, sumDistinct
# Select all from the temporary view
# Convert the column to an integer with the new column name
sqlDF.select(sum("<new column name>")) # Select sum
sqlDF.select(sumDistinct("<new column name>")) # Select sum
distinct
```

```
+-----+
|sum(reactions)|
+-----+
|          8382|
+-----+


+-----+
|sum(DISTINCT reactions)|
+-----+
|          5557|
+-----+
```

Aggregation Functions (Chambers and Zaharia, 2018)

- avg (average) is used for calculating an average value of a column in a Dataframe.

```
from pyspark.sql.functions import avg
# Select all from temporary view
# Convert the column to integer with the new column name
sqlDF.select(avg("<new column name>")) # Select average
```

Joins



Joins (Chambers and Zaharia, 2018)

- Join combines two sets of data based on the key of the left and right datasets (tables).
- Spark discards the unmatched rows and then returns matched rows.
- Join evaluates the result using a join expression.

```
join_column = sqlDF1["<column name>"] == \  
                sqlDF2["<column name>"] # Identify column  
                                         # for matching
```


Joins (Chambers and Zaharia, 2018)

- Inner join keeps rows with keys matched in the left and right datasets.

```
sqlDF1.join(sqlDF2, join_column).show() # Inner join and  
                                         # show results
```

```
joinType = "inner"  
sqlDF1.join(sqlDF2, join_column, joinType).show() # Inner  
                                                    # join and show results  
                                                    # with using joinType
```

Joins (Chambers and Zaharia, 2018)

- Inner join example result:

Table	status_id	status_type	status_published	num_reactions	num_comments	Table	status_id	status_type	status_published	num_reactions	num_shares
FB2	246675545449582_1...	video	4/22/2018 6:00	529	512	FB3	246675545449582_1...	video	4/22/2018 6:00	529	262
FB2	246675545449582_1...	photo	4/21/2018 22:45	150	0	FB3	246675545449582_1...	photo	4/21/2018 22:45	150	0
FB2	246675545449582_1...	video	4/21/2018 6:17	227	236	FB3	246675545449582_1...	video	4/21/2018 6:17	227	57
FB2	246675545449582_1...	photo	4/21/2018 2:29	111	0	FB3	246675545449582_1...	photo	4/21/2018 2:29	111	0
FB2	246675545449582_1...	photo	4/18/2018 3:22	213	0	FB3	246675545449582_1...	photo	4/18/2018 3:22	213	0
FB2	246675545449582_1...	photo	4/18/2018 2:14	217	6	FB3	246675545449582_1...	photo	4/18/2018 2:14	217	0
FB2	246675545449582_1...	video	4/18/2018 0:24	503	614	FB3	246675545449582_1...	video	4/18/2018 0:24	503	72
FB2	246675545449582_1...	video	4/17/2018 7:42	295	453	FB3	246675545449582_1...	video	4/17/2018 7:42	295	53
FB2	246675545449582_1...	photo	4/17/2018 3:33	203	1	FB3	246675545449582_1...	photo	4/17/2018 3:33	203	0

Joins (Chambers and Zaharia, 2018)

- Outer join keeps rows with keys in either the left or right datasets.

joinType = "outer"

```
sqlDF1.join(sqlDF2, join_column, joinType).show()
```

Table	status_id	status_type	status_published	num_reactions	num_comments	Table	status_id	status_type	status_published	num_reactions	num_shares
FB2	246675545449582_1...	photo	3/12/2018 5:51	145	9	null	null	null	null	null	null
FB2	246675545449582_1...	video	3/17/2018 7:47	90	78	null	null	null	null	null	null
FB2	246675545449582_1...	video	3/17/2018 8:07	75	36	null	null	null	null	null	null
FB2	246675545449582_1...	video	3/19/2018 22:34	37	0	null	null	null	null	null	null
FB2	246675545449582_1...	photo	3/20/2018 0:15	102	0	null	null	null	null	null	null
FB2	246675545449582_1...	video	3/20/2018 1:28	18	0	null	null	null	null	null	null
FB2	246675545449582_1...	photo	3/20/2018 1:54	98	0	null	null	null	null	null	null
FB2	246675545449582_1...	photo	3/21/2018 7:46	227	7	null	null	null	null	null	null
FB2	246675545449582_1...	photo	3/21/2018 8:40	234	15	null	null	null	null	null	null
FB2	246675545449582_1...	photo	3/22/2018 1:25	152	2	null	null	null	null	null	null
null	null	null	null	null	null	FB3	246675545449582_1...	video	3/23/2018 7:09	221	36
null	null	null	null	null	null	FB3	246675545449582_1...	video	3/26/2018 8:28	150	47
null	null	null	null	null	null	FB3	246675545449582_1...	video	3/30/2018 8:28	135	79
null	null	null	null	null	null	FB3	246675545449582_1...	video	4/1/2018 5:16	332	30
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/5/2018 9:23	346	0
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/8/2018 2:23	209	0
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/8/2018 5:10	313	2
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/9/2018 2:06	222	0
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/10/2018 1:01	210	3
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/11/2018 4:53	170	1

Joins (Chambers and Zaharia, 2018)

- Left outer join keeps rows with keys in the left dataset.
- Right outer join keeps rows with keys in the right dataset.

```
joinTypeLeft = "left_outer" # Left outer join
sqlDF1.join(sqlDF2, join_column, joinTypeLeft).show()
joinTypeRight = "right_outer" # Right outer join
sqlDF1.join(sqlDF2, join_column, joinTypeRight).show()
```

Left Outer Join Example Result


Table	status_id	status_type	status_published	num_reactions	num_comments	Table	status_id	status_type	status_published	num_reactions	num_shares
FB2	246675545449582_1...	video	4/22/2018 6:00	529	512	FB3	246675545449582_1...	video	4/22/2018 6:00	529	262
FB2	246675545449582_1...	photo	4/21/2018 22:45	150	0	FB3	246675545449582_1...	photo	4/21/2018 22:45	150	0
FB2	246675545449582_1...	video	4/21/2018 6:17	227	236	FB3	246675545449582_1...	video	4/21/2018 6:17	227	57
FB2	246675545449582_1...	photo	4/21/2018 2:29	111	0	FB3	246675545449582_1...	photo	4/21/2018 2:29	111	0
FB2	246675545449582_1...	photo	4/18/2018 3:22	213	0	FB3	246675545449582_1...	photo	4/18/2018 3:22	213	0
FB2	246675545449582_1...	photo	4/18/2018 2:14	217	6	FB3	246675545449582_1...	photo	4/18/2018 2:14	217	0
FB2	246675545449582_1...	video	4/18/2018 0:24	503	614	FB3	246675545449582_1...	video	4/18/2018 0:24	503	72
FB2	246675545449582_1...	video	4/17/2018 7:42	295	453	FB3	246675545449582_1...	video	4/17/2018 7:42	295	53
FB2	246675545449582_1...	photo	4/17/2018 3:33	203	1	FB3	246675545449582_1...	photo	4/17/2018 3:33	203	0
FB2	246675545449582_1...	photo	3/22/2018 1:25	152	2	null	null	null	null	null	null
FB2	246675545449582_1...	photo	3/21/2018 8:40	234	15	null	null	null	null	null	null
FB2	246675545449582_1...	photo	3/21/2018 7:46	227	7	null	null	null	null	null	null
FB2	246675545449582_1...	photo	3/20/2018 1:54	98	0	null	null	null	null	null	null
FB2	246675545449582_1...	video	3/20/2018 1:28	18	0	null	null	null	null	null	null
FB2	246675545449582_1...	photo	3/20/2018 0:15	102	0	null	null	null	null	null	null
FB2	246675545449582_1...	video	3/19/2018 22:34	37	0	null	null	null	null	null	null
FB2	246675545449582_1...	video	3/17/2018 8:07	75	36	null	null	null	null	null	null
FB2	246675545449582_1...	video	3/17/2018 7:47	90	78	null	null	null	null	null	null
FB2	246675545449582_1...	photo	3/12/2018 5:51	145	9	null	null	null	null	null	null

Right Outer Join Example Result

Table	status_id	status_type	status_published	num_reactions	num_comments	Table	status_id	status_type	status_published	num_reactions	num_shares
FB2	246675545449582_1...	video	4/22/2018 6:00	529	512	FB3	246675545449582_1...	video	4/22/2018 6:00	529	262
FB2	246675545449582_1...	photo	4/21/2018 22:45	150	0	FB3	246675545449582_1...	photo	4/21/2018 22:45	150	0
FB2	246675545449582_1...	video	4/21/2018 6:17	227	236	FB3	246675545449582_1...	video	4/21/2018 6:17	227	57
FB2	246675545449582_1...	photo	4/21/2018 2:29	111	0	FB3	246675545449582_1...	photo	4/21/2018 2:29	111	0
FB2	246675545449582_1...	photo	4/18/2018 3:22	213	0	FB3	246675545449582_1...	photo	4/18/2018 3:22	213	0
FB2	246675545449582_1...	photo	4/18/2018 2:14	217	6	FB3	246675545449582_1...	photo	4/18/2018 2:14	217	0
FB2	246675545449582_1...	video	4/18/2018 0:24	503	614	FB3	246675545449582_1...	video	4/18/2018 0:24	503	72
FB2	246675545449582_1...	video	4/17/2018 7:42	295	453	FB3	246675545449582_1...	video	4/17/2018 7:42	295	53
FB2	246675545449582_1...	photo	4/17/2018 3:33	203	1	FB3	246675545449582_1...	photo	4/17/2018 3:33	203	0
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/11/2018 4:53	170	1
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/10/2018 1:01	210	3
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/9/2018 2:06	222	0
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/8/2018 5:10	313	2
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/8/2018 2:23	209	0
null	null	null	null	null	null	FB3	246675545449582_1...	photo	4/5/2018 9:23	346	0
null	null	null	null	null	null	FB3	246675545449582_1...	video	4/1/2018 5:16	332	30
null	null	null	null	null	null	FB3	246675545449582_1...	video	3/30/2018 8:28	135	79
null	null	null	null	null	null	FB3	246675545449582_1...	video	3/26/2018 8:28	150	47
null	null	null	null	null	null	FB3	246675545449582_1...	video	3/23/2018 7:09	221	36

Assignment (1 point)

- Please implement the codes in slides 9 to 28 in one file and show the results to get 1 point.





References

- Chambers, B., & Zaharia, M. (2018). *Spark: The definitive guide: Big data processing made simple*. "O'Reilly Media, Inc."
 - Antolínez García, A. (2023). *Hands-on Guide to Apache Spark 3: Build Scalable Computing Engines for Batch and Stream Data Processing*. Berkeley, CA: Apress.
- 